

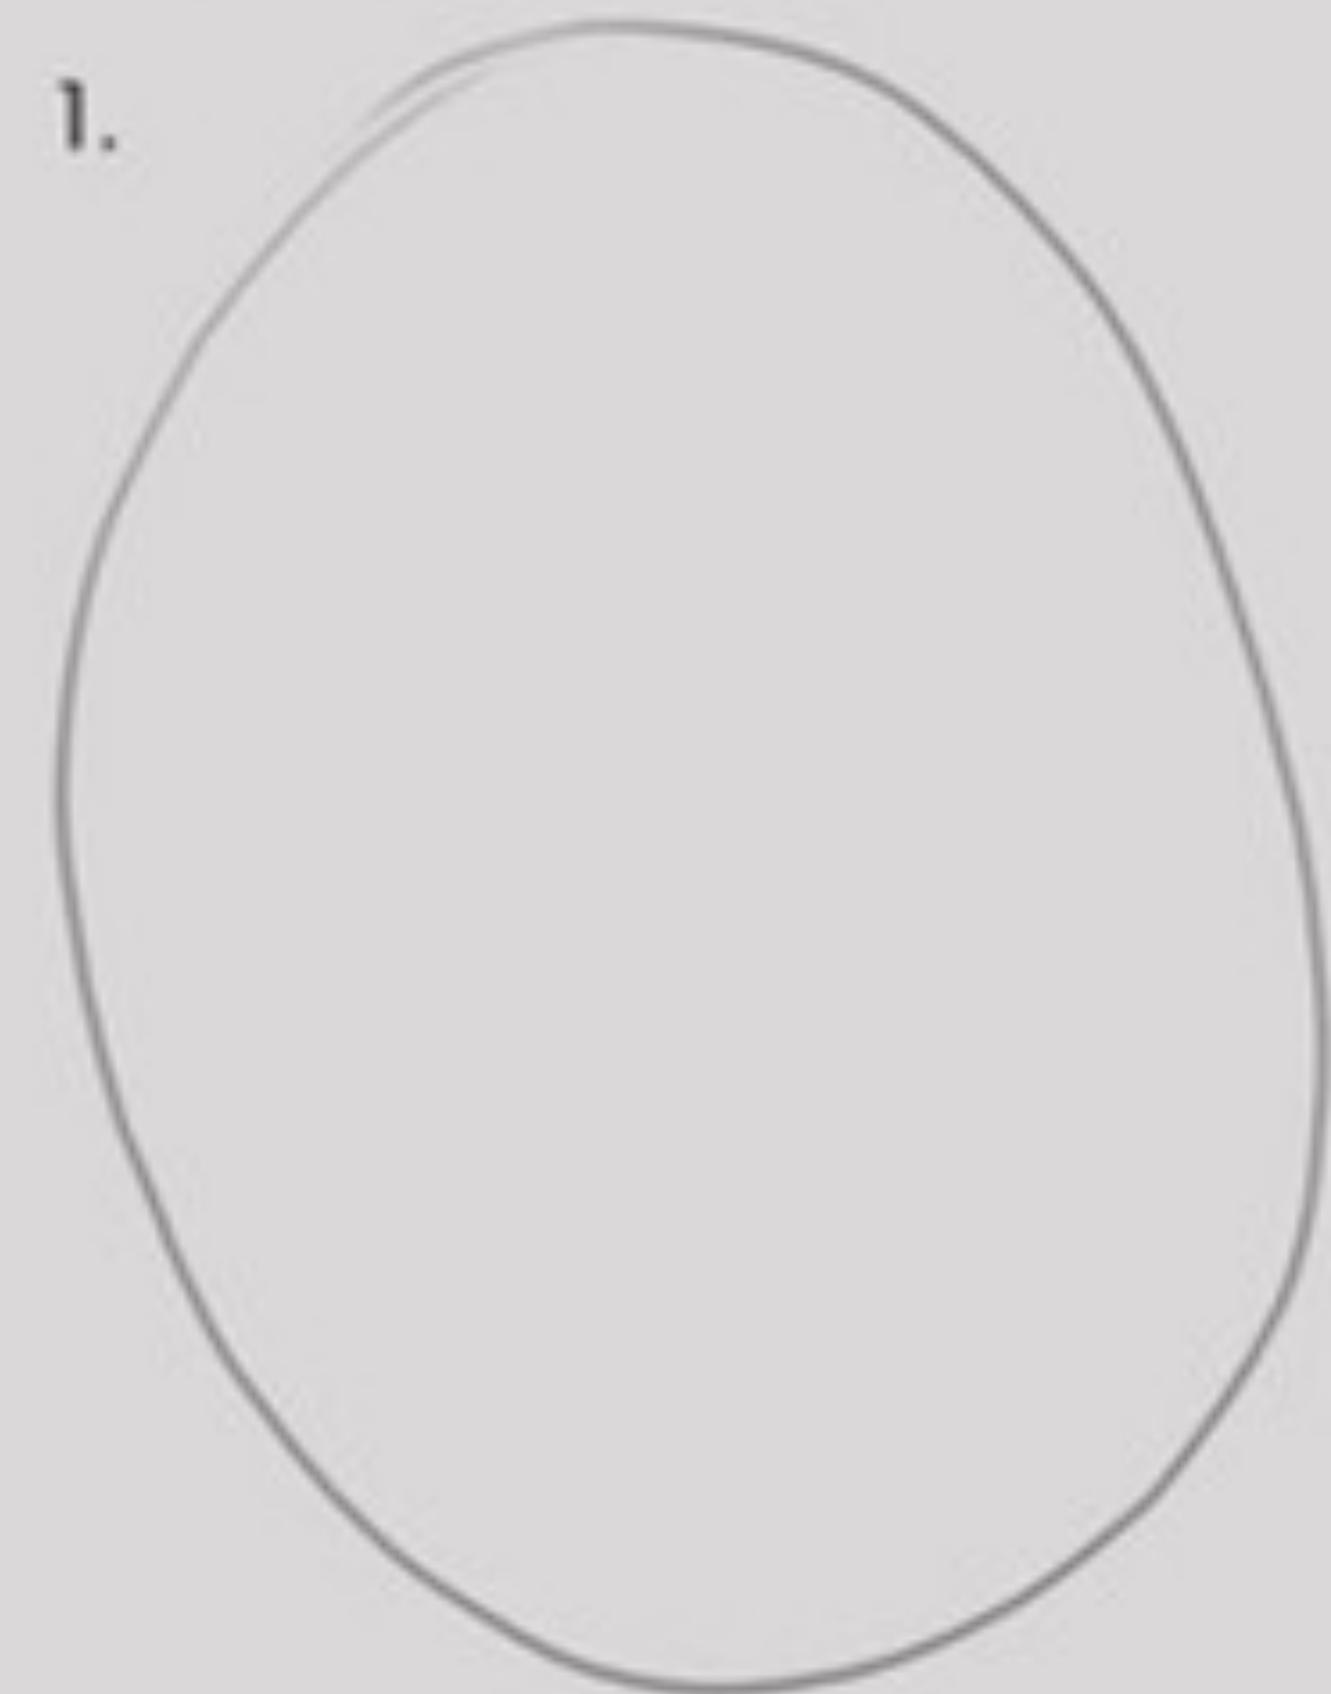
Coding papers in PYTORCH (or any other framework)

Soumith Chintala



How to draw a head

1.



1. Draw a circle

2.



2. Draw the rest of the head



Let's discuss

- Unsupervised representation learning with deep convolutional generative adversarial networks

Radford et. al.

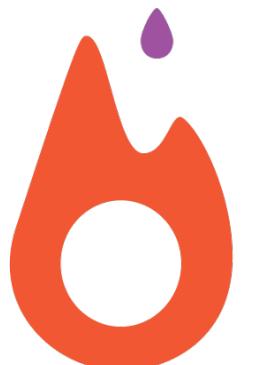


Start with the experiments section

4 DETAILS OF ADVERSARIAL TRAINING

We trained DCGANs on three datasets, Large-scale Scene Understanding (LSUN) (Yu et al., 2015), Imagenet-1k and a newly assembled Faces dataset. Details on the usage of each of these datasets are given below.

No pre-processing was applied to training images besides scaling to the range of the tanh activation function [-1, 1]. All models were trained with mini-batch stochastic gradient descent (SGD) with a mini-batch size of 128. All weights were initialized from a zero-centered Normal distribution with standard deviation 0.02. In the LeakyReLU, the slope of the leak was set to 0.2 in all models. While previous GAN work has used momentum to accelerate training, we used the Adam optimizer (Kingma & Ba, 2014) with tuned hyperparameters. We found the suggested learning rate of 0.001, to be too high, using 0.0002 instead. Additionally, we found leaving the momentum term β_1 at the



Start with the experiments section

- Helps make a list of tasks



Start with the experiments section

- Which datasets were used?



Start with the experiments section

- Which datasets were used?
- architecture of the neural network (or model)



Start with the experiments section

- Which datasets were used?
- architecture of the neural network (or model)
- What is the loss function?



Start with the experiments section

- Which datasets were used?
- architecture of the neural network (or model)
- What is the loss function?
If it is a standard function, only one or two words are said about it. For example: "Cross-entropy loss is used"



Find Algorithm 1

- Some papers give easier to understand pseudo-code

Algorithm 1 WGAN with gradient penalty. We use default values of $\lambda = 10$, $n_{\text{critic}} = 5$, $\alpha = 0.0001$, $\beta_1 = 0$, $\beta_2 = 0.9$.

Require: The gradient penalty coefficient λ , the number of critic iterations per generator iteration n_{critic} , the batch size m , Adam hyperparameters α, β_1, β_2 .

Require: initial critic parameters w_0 , initial generator parameters θ_0 .

```
1: while  $\theta$  has not converged do
2:   for  $t = 1, \dots, n_{\text{critic}}$  do
3:     for  $i = 1, \dots, m$  do
4:       Sample real data  $\mathbf{x} \sim \mathbb{P}_r$ , latent variable  $\mathbf{z} \sim p(\mathbf{z})$ , a random number  $\epsilon \sim U[0, 1]$ .
5:        $\tilde{\mathbf{x}} \leftarrow G_\theta(\mathbf{z})$ 
6:        $\hat{\mathbf{x}} \leftarrow \epsilon \mathbf{x} + (1 - \epsilon) \tilde{\mathbf{x}}$ 
7:        $L^{(i)} \leftarrow D_w(\tilde{\mathbf{x}}) - D_w(\mathbf{x}) + \lambda(\|\nabla_{\hat{\mathbf{x}}} D_w(\hat{\mathbf{x}})\|_2 - 1)^2$ 
8:     end for
9:      $w \leftarrow \text{Adam}(\nabla_w \frac{1}{m} \sum_{i=1}^m L^{(i)}, w, \alpha, \beta_1, \beta_2)$ 
10:   end for
11:   Sample a batch of latent variables  $\{\mathbf{z}^{(i)}\}_{i=1}^m \sim p(\mathbf{z})$ .
12:    $\theta \leftarrow \text{Adam}(\nabla_\theta \frac{1}{m} \sum_{i=1}^m -D_w(G_\theta(\mathbf{z})), \theta, \alpha, \beta_1, \beta_2)$ 
13: end while
```



Write and verify dataset loader

- Before starting paper implementation!!!
-



Write and verify dataset loader

- Before starting paper implementation!!!
- Verify visually and with simple unit tests



Write and verify dataset loader

- Before starting paper implementation!!!
- Verify visually and with simple unit tests



Write and verify dataset loader

- Before starting paper implementation!!!
- Verify visually and with simple unit tests
- if `__name__ == "__main__"` is useful



Write and verify dataset loader

- Before starting paper implementation!!!
- Verify visually and with simple unit tests
- if `__name__ == "__main__"` is useful

```
126  if __name__ == '__main__':
127      #lsun = LSUNClassDataset(db_path='/home/soumith/local/lsun/train/bedroom_train_lmdb')
128      #a = lsun[0]
129      lsun = LSUNDataset(db_path='/home/soumith/local/lsun/train',
130                          classes=['bedroom_train', 'church_outdoor_train'])
131      print(lsun.classes)
132      print(lsun.dbs)
133      a, t = lsun[len(lsun)-1]
134      print(a)
135      print(t)
```



Starting with the paper

- Aim for the smallest dataset first



Starting with the paper

- Aim for the smallest dataset first
quick iteration, faster initial development



Starting with the paper

- Aim for the smallest dataset first
quick iteration, faster initial development
CIFAR10 is a good sanity check. MNIST is not!



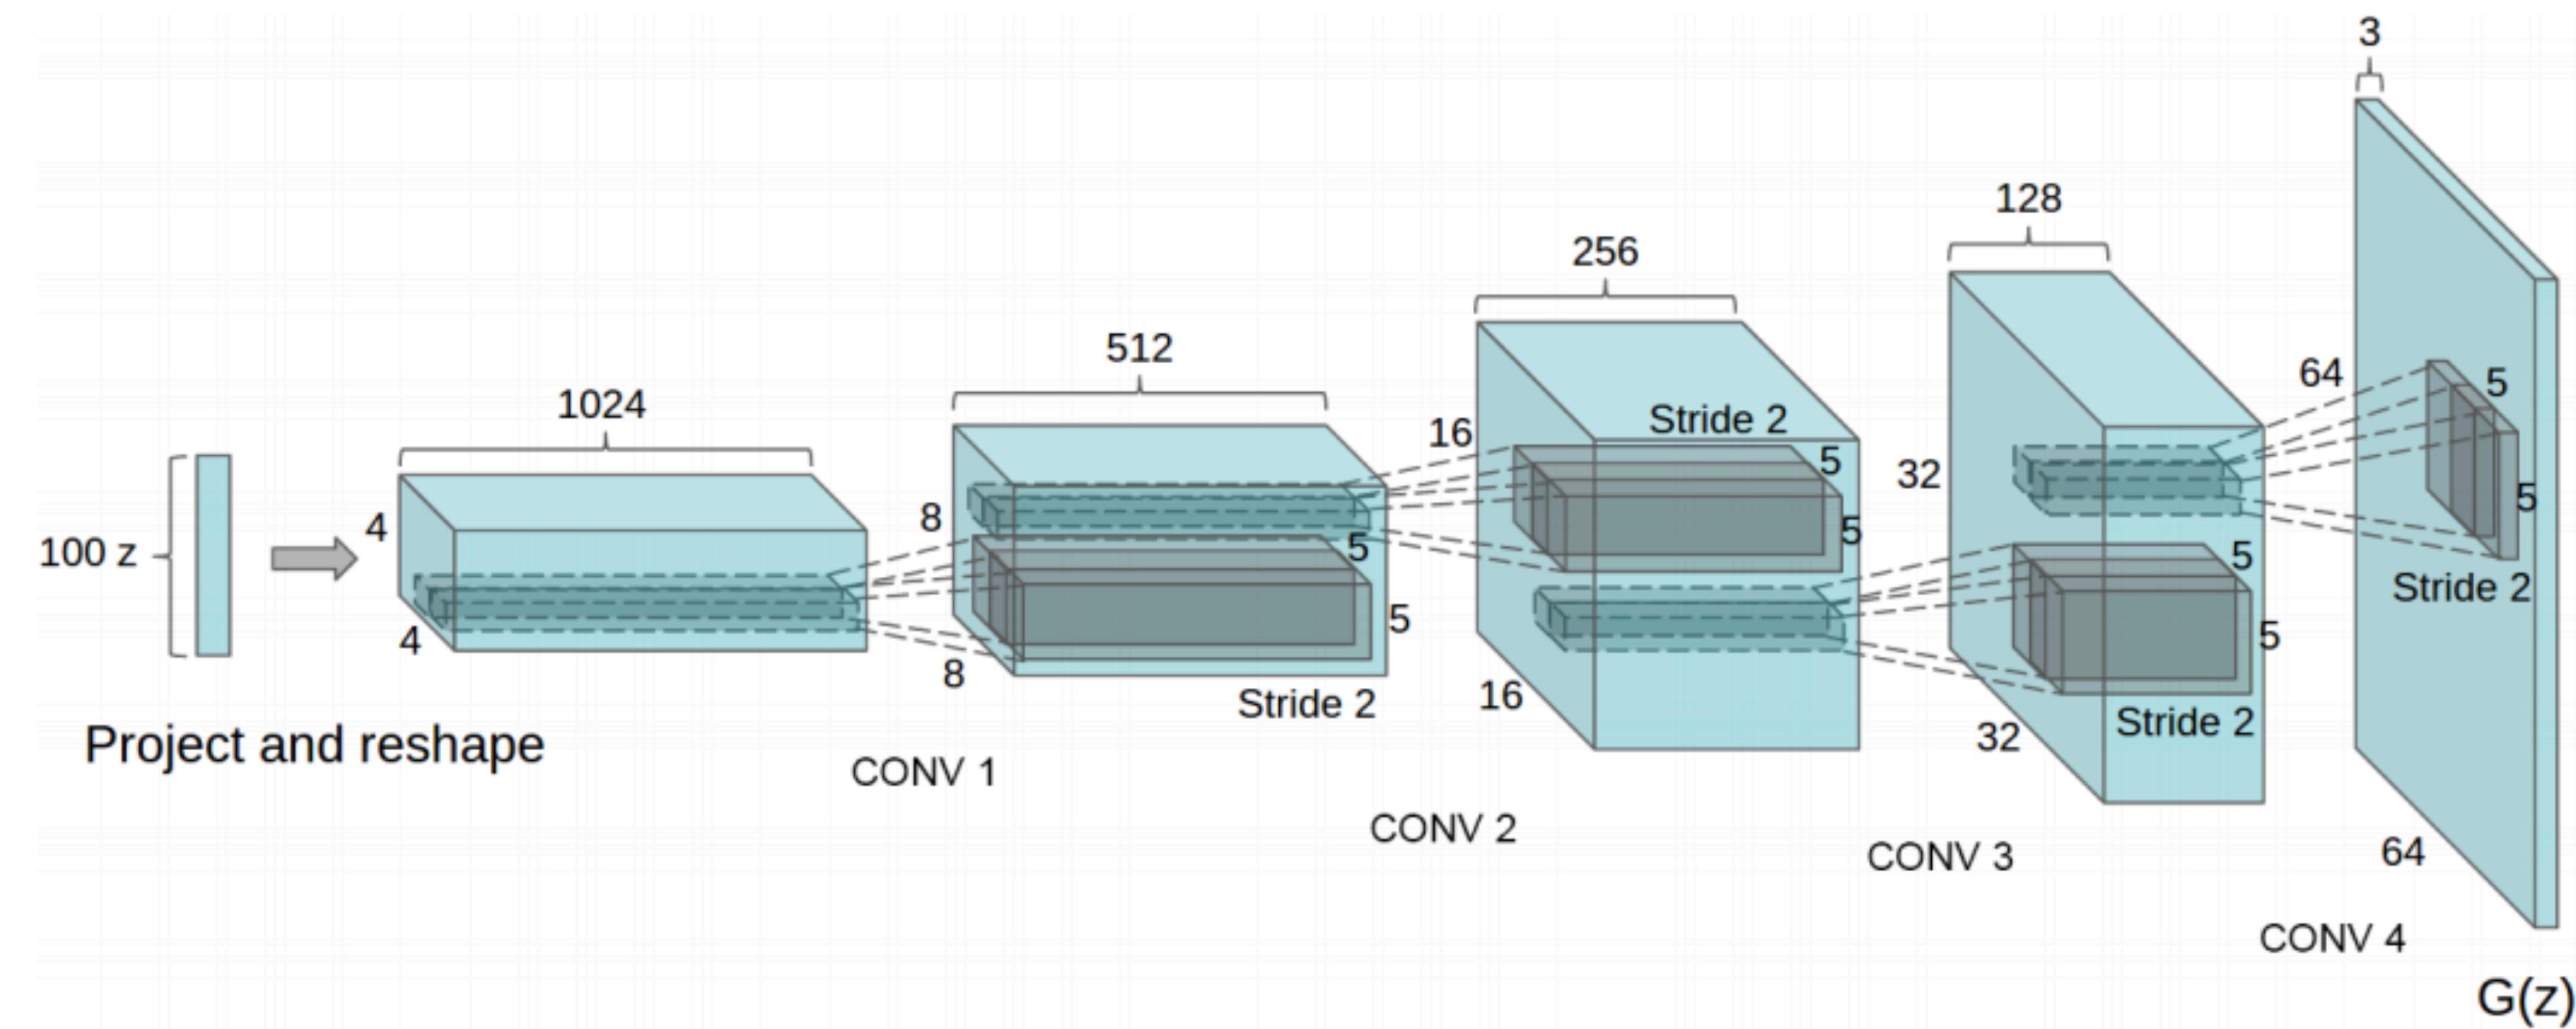
Starting with the paper

- Aim for the smallest dataset first
- Implement the model



Model definition

- Aim for the smallest dataset first
- Implement the model



Model definition

- Aim for the smallest dataset first
- Implement the model

Clear model details not available in most papers (sadly!)



Model definition

- Aim for the smallest dataset first
- Implement the model

Clear model details not available in most papers (sadly!)

See appendix



Model definition

- Aim for the smallest dataset first
- Implement the model

Clear model details not available in most papers (sadly!)

See appendix

Sometimes, model definition is not easy:



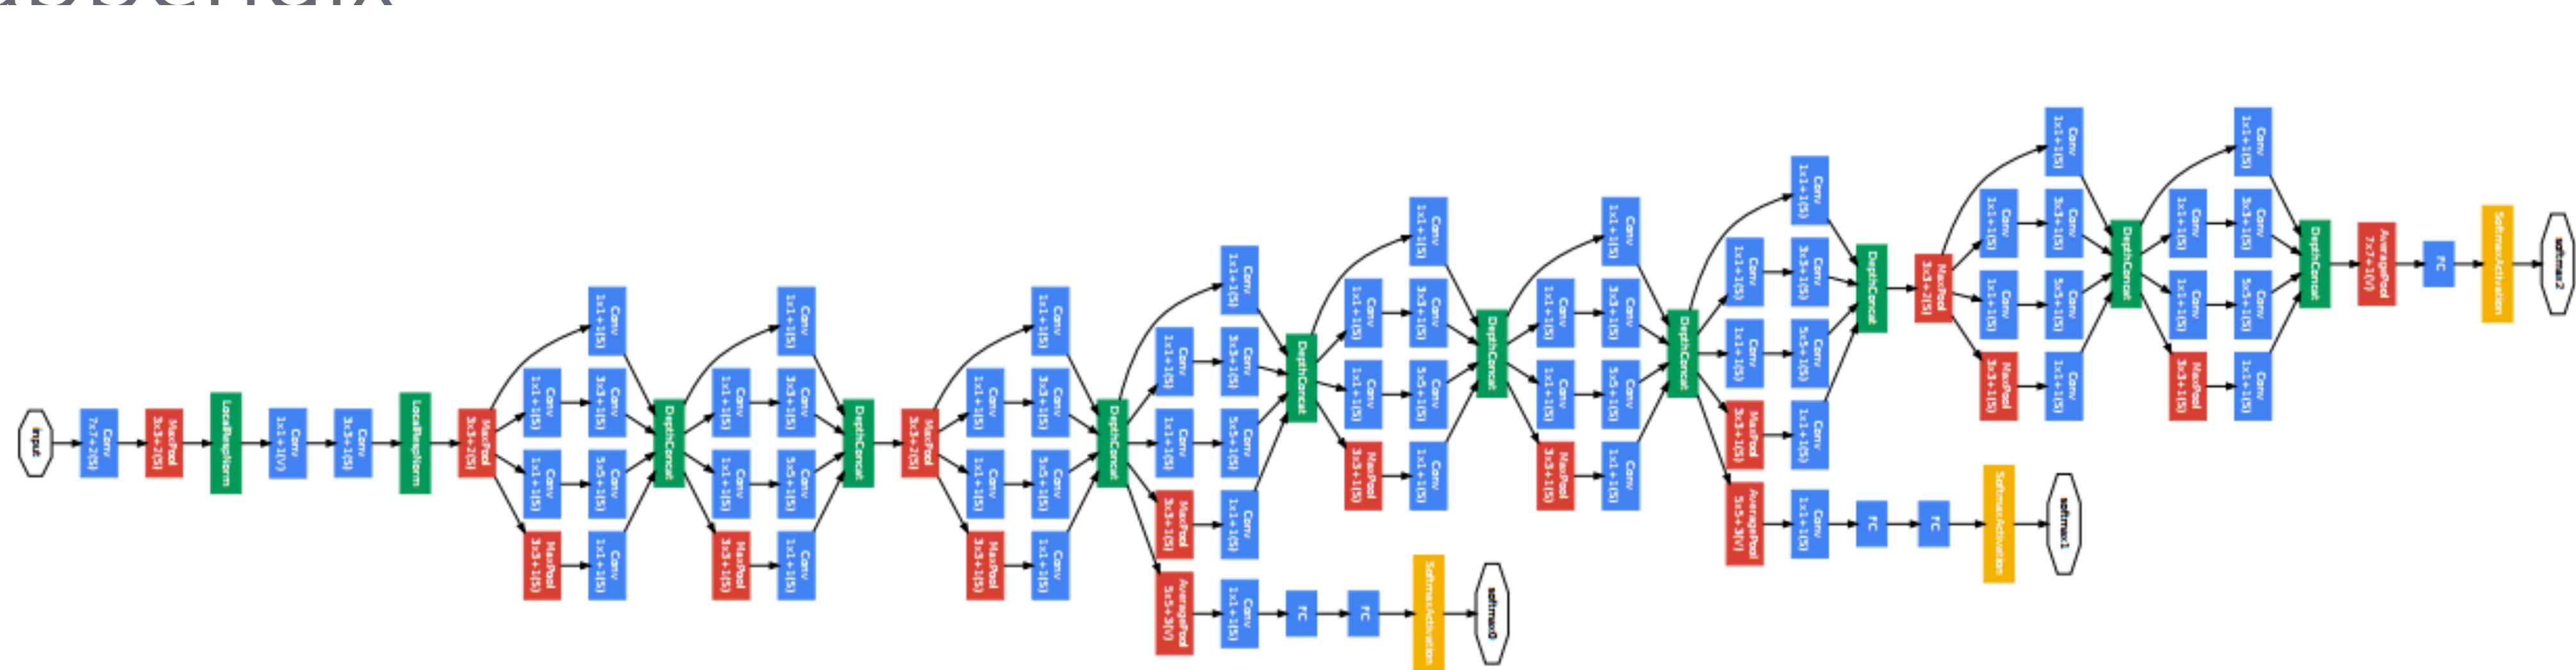
Model definition

- Aim for the smallest dataset first
- Implement the model

Clear model details not available in most papers (sadly!)

See appendix

Son



Model definition

- Aim for the smallest dataset first
- Implement the model

Clear model details not available in most papers (sadly!)

See appendix

Sometimes, model definition is not easy:



Model def

- Aim for the sr
- Implement the

Clear model c

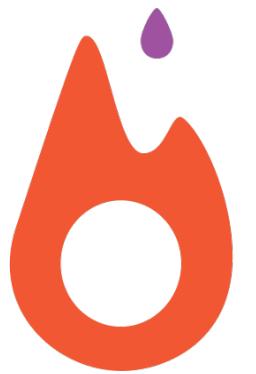
See appendix

Sometimes, m

```
class _netG(nn.Module):
    def __init__(self, ngpu):
        super(_netG, self).__init__()
        self.ngpu = ngpu
        self.main = nn.Sequential(
            # input is Z, going into a convolution
            nn.ConvTranspose2d(      nz, ngf * 8, 4, 1, 0, bias=False),
            nn.BatchNorm2d(ngf * 8),
            nn.ReLU(True),
            # state size. (ngf*8) x 4 x 4
            nn.ConvTranspose2d(ngf * 8, ngf * 4, 4, 2, 1, bias=False),
            nn.BatchNorm2d(ngf * 4),
            nn.ReLU(True),
            # state size. (ngf*4) x 8 x 8
            nn.ConvTranspose2d(ngf * 4, ngf * 2, 4, 2, 1, bias=False),
            nn.BatchNorm2d(ngf * 2),
            nn.ReLU(True),
            # state size. (ngf*2) x 16 x 16
            nn.ConvTranspose2d(ngf * 2,      ngf, 4, 2, 1, bias=False),
            nn.BatchNorm2d(ngf),
            nn.ReLU(True),
            # state size. (ngf) x 32 x 32
            nn.ConvTranspose2d(      ngf,      nc, 4, 2, 1, bias=False),
            nn.Tanh()
            # state size. (nc) x 64 x 64
        )

    def forward(self, input):
        if isinstance(input.data, torch.cuda.FloatTensor) and self.ngpu > 1:
            output = nn.parallel.data_parallel(self.main, input, range(self.ngpu))
        else:
            output = self.main(input)
        return output
```

s (sadly!)



Model def

- Aim for the sr
 - Implement th
- Clear model c
See appendix
Sometimes, n

```
class _netD(nn.Module):  
    def __init__(self, ngpu):  
        super(_netD, self).__init__()  
        self.ngpu = ngpu  
        self.main = nn.Sequential(  
            # input is (nc) x 64 x 64  
            nn.Conv2d(nc, ndf, 4, 2, 1, bias=False),  
            nn.LeakyReLU(0.2, inplace=True),  
            # state size. (ndf) x 32 x 32  
            nn.Conv2d(ndf, ndf * 2, 4, 2, 1, bias=False),  
            nn.BatchNorm2d(ndf * 2),  
            nn.LeakyReLU(0.2, inplace=True),  
            # state size. (ndf*2) x 16 x 16  
            nn.Conv2d(ndf * 2, ndf * 4, 4, 2, 1, bias=False),  
            nn.BatchNorm2d(ndf * 4),  
            nn.LeakyReLU(0.2, inplace=True),  
            # state size. (ndf*4) x 8 x 8  
            nn.Conv2d(ndf * 4, ndf * 8, 4, 2, 1, bias=False),  
            nn.BatchNorm2d(ndf * 8),  
            nn.LeakyReLU(0.2, inplace=True),  
            # state size. (ndf*8) x 4 x 4  
            nn.Conv2d(ndf * 8, 1, 4, 1, 0, bias=False),  
            nn.Sigmoid()  
        )  
  
        def forward(self, input):  
            if isinstance(input.data, torch.cuda.FloatTensor) and self.ngpu > 1:  
                output = nn.parallel.data_parallel(self.main, input, range(self.ngpu))  
            else:  
                output = self.main(input)  
  
            return output.view(-1, 1).squeeze(1)
```

; (sadly!)



Model def

- Aim for the sr
- Implement th

Clear model c

WEIGHT INITIALIZATION IS IMPORTANT

See appendix

Sometimes, n

```
class _netD(nn.Module):
    def __init__(self, ngpu):
        super(_netD, self).__init__()
        self.ngpu = ngpu
        self.main = nn.Sequential(
            # input is (nc) x 64 x 64
            nn.Conv2d(nc, ndf, 4, 2, 1, bias=False),
            nn.LeakyReLU(0.2, inplace=True),
            # state size. (ndf) x 32 x 32
            nn.Conv2d(ndf, ndf * 2, 4, 2, 1, bias=False),
            nn.BatchNorm2d(ndf * 2),
            nn.LeakyReLU(0.2, inplace=True),
            # state size. (ndf*2) x 16 x 16
            nn.Conv2d(ndf * 2, ndf * 4, 4, 2, 1, bias=False),
            nn.BatchNorm2d(ndf * 4),
            nn.LeakyReLU(0.2, inplace=True),
            # state size. (ndf*8) x 4 x 4
            nn.Conv2d(ndf * 4, ndf * 8, 4, 2, 1, bias=False),
            nn.BatchNorm2d(ndf * 8),
            nn.LeakyReLU(0.2, inplace=True),
            # state size. (ndf*8) x 4 x 4
            nn.Conv2d(ndf * 8, 1, 4, 1, 0, bias=False),
            nn.Sigmoid()
        )

    def forward(self, input):
        if isinstance(input.data, torch.cuda.FloatTensor) and self.ngpu > 1:
            output = nn.parallel.data_parallel(self.main, input, range(self.ngpu))
        else:
            output = self.main(input)

        return output.view(-1, 1).squeeze(1)
```



Training loop and optimizer

Find the correct hyper parameters from paper

- Hidden / hard-to-find in most papers



Training loop and optimizer

Find the correct hyper parameters from paper

- Hidden / hard-to-find in most papers

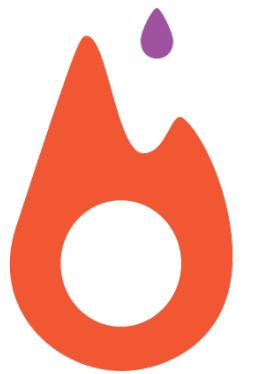
No pre-processing was applied to training images besides scaling to the range of the tanh activation function [-1, 1]. All models were trained with mini-batch stochastic gradient descent (SGD) with a mini-batch size of 128. All weights were initialized from a zero-centered Normal distribution with standard deviation 0.02. In the LeakyReLU, the slope of the leak was set to 0.2 in all models. While previous GAN work has used momentum to accelerate training, we used the Adam optimizer (Kingma & Ba, 2014) with tuned hyperparameters. We found the suggested learning rate of 0.001, to be too high, using 0.0002 instead. Additionally, we found leaving the momentum term β_1 at the suggested value of 0.9 resulted in training oscillation and instability while reducing it to 0.5 helped stabilize training.

```
# setup optimizer
optimizerD = optim.Adam(netD.parameters(), lr=opt.lr, betas=(opt.beta1, 0.999))
optimizerG = optim.Adam(netG.parameters(), lr=opt.lr, betas=(opt.beta1, 0.999))
```



Training loop and optimizer

```
for epoch in range(opt.niter):  
    for i, data in enumerate(dataloader, 0):
```



Training loop and optimizer

```
for epoch in range(opt.niter):
    for i, data in enumerate(dataloader, 0):
        #####
        # (1) Update D network: maximize log(D(x)) + log(1 - D(G(z)))
        #####
        # train with real
        netD.zero_grad()
        real_cpu, _ = data
        batch_size = real_cpu.size(0)
        if opt.cuda:
            real_cpu = real_cpu.cuda()
        inputv.resize_as_(real_cpu).copy_(real_cpu)
        labelv.resize_(batch_size).fill_(real_label)
        inputv = Variable(inputv)
        labelv = Variable(labelv)

        output = netD(inputv)
        errD_real = criterion(output, labelv)
        errD_real.backward()
        D_x = output.data.mean()

        # train with fake
        noisev.resize_(batch_size, nz, 1, 1).normal_(0, 1)
        noisev = Variable(noisev)
        fakev = netG(noisev)
        labelv = Variable(labelv.fill_(fake_label))
        output = netD(fakev.detach())
        errD_fake = criterion(output, labelv)
        errD_fake.backward()
        D_G_z1 = output.data.mean()
        errD = errD_real + errD_fake
        optimizerD.step()
```

Training loop and optimizer

- Print statistics



Training loop and optimizer

- Print statistics
- Save checkpoints



Training loop and optimizer

- Print statistics
- Save checkpoints
- profile!

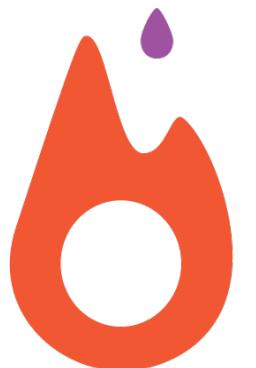


Finally!



Finally!

```
1 from __future__ import print_function
2 import argparse
3 import os
4 import random
5 import torch
6 import torch.nn as nn
7 import torch.nn.parallel
8 import torch.backends.cudnn as cudnn
9 import torch.optim as optim
10 import torch.utils.data
11 import torchvision.datasets as dset
12 import torchvision.transforms as transforms
13 import torchvision.utils as vutils
14 from torch.autograd import Variable
15 --
```



Finally!

```
17 parser = argparse.ArgumentParser()
18 parser.add_argument('--dataset', required=True, help='cifar10 | lsun | imagenet | folder | lfw | fake')
19 parser.add_argument('--dataroot', required=True, help='path to dataset')
20 parser.add_argument('--workers', type=int, help='number of data loading workers', default=2)
21 parser.add_argument('--batchSize', type=int, default=64, help='input batch size')
22 parser.add_argument('--imageSize', type=int, default=64, help='the height / width of the input image to network')
23 parser.add_argument('--nz', type=int, default=100, help='size of the latent z vector')
24 parser.add_argument('--ngf', type=int, default=64)
25 parser.add_argument('--ndf', type=int, default=64)
26 parser.add_argument('--niter', type=int, default=25, help='number of epochs to train for')
27 parser.add_argument('--lr', type=float, default=0.0002, help='learning rate, default=0.0002')
28 parser.add_argument('--beta1', type=float, default=0.5, help='beta1 for adam. default=0.5')
29 parser.add_argument('--cuda', action='store_true', help='enables cuda')
30 parser.add_argument('--ngpu', type=int, default=1, help='number of GPUs to use')
31 parser.add_argument('--netG', default='', help="path to netG (to continue training)")
32 parser.add_argument('--netD', default='', help="path to netD (to continue training)")
33 parser.add_argument('--outf', default='.', help='folder to output images and model checkpoints')
34 parser.add_argument('--manualSeed', type=int, help='manual seed')
35
```



Finally!

```
36     opt = parser.parse_args()
37     print(opt)
38
39     try:
40         os.makedirs(opt.outf)
41     except OSError:
42         pass
43
44     if opt.manualSeed is None:
45         opt.manualSeed = random.randint(1, 10000)
46     print("Random Seed: ", opt.manualSeed)
47     random.seed(opt.manualSeed)
48     torch.manual_seed(opt.manualSeed)
49     if opt.cuda:
50         torch.cuda.manual_seed_all(opt.manualSeed)
51
52     cudnn.benchmark = True
53
54     if torch.cuda.is_available() and not opt.cuda:
55         print("WARNING: You have a CUDA device, so you should probably run with --cuda")
56
```



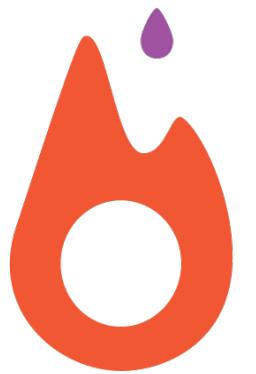
Finally

```
57 if opt.dataset in ['imagenet', 'folder', 'lfw']:
58     # folder dataset
59     dataset = dset.ImageFolder(root=opt.dataroot,
60                               transform=transforms.Compose([
61                                 transforms.Resize(opt.imageSize),
62                                 transforms.CenterCrop(opt.imageSize),
63                                 transforms.ToTensor(),
64                                 transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5)),
65                               ]))
66 elif opt.dataset == 'lsun':
67     dataset = dset.LSUN(db_path=opt.dataroot, classes=['bedroom_train'],
68                         transform=transforms.Compose([
69                           transforms.Resize(opt.imageSize),
70                           transforms.CenterCrop(opt.imageSize),
71                           transforms.ToTensor(),
72                           transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5)),
73                         ]))
74 elif opt.dataset == 'cifar10':
75     dataset = dset.CIFAR10(root=opt.dataroot, download=True,
76                           transform=transforms.Compose([
77                             transforms.Resize(opt.imageSize),
78                             transforms.ToTensor(),
79                             transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5)),
80                           ]))
81 elif opt.dataset == 'fake':
82     dataset = dset.FakeData(image_size=(3, opt.imageSize, opt.imageSize),
83                             transform=transforms.ToTensor())
84 assert dataset
85 dataloader = torch.utils.data.DataLoader(dataset, batch_size=opt.batchSize,
86                                         shuffle=True, num_workers=int(opt.workers))
```



Finally!

```
88    ngpu = int(opt.ngpu)
89    nz = int(opt.nz)
90    ngf = int(opt.ngf)
91    ndf = int(opt.ndf)
92    nc = 3
93
94
95    # custom weights initialization called on netG and netD
96    def weights_init(m):
97        classname = m.__class__.__name__
98        if classname.find('Conv') != -1:
99            m.weight.data.normal_(0.0, 0.02)
100       elif classname.find('BatchNorm') != -1:
101           m.weight.data.normal_(1.0, 0.02)
102           m.bias.data.fill_(0)
103
```



Finally!

```
105 class _netG(nn.Module):
106     def __init__(self, ngpu):
107         super(_netG, self).__init__()
108         self.ngpu = ngpu
109         self.main = nn.Sequential(
110             # input is Z, going into a convolution
111             nn.ConvTranspose2d(      nz, ngf * 8, 4, 1, 0, bias=False),
112             nn.BatchNorm2d(ngf * 8),
113             nn.ReLU(True),
114             # state size. (ngf*8) x 4 x 4
115             nn.ConvTranspose2d(ngf * 8, ngf * 4, 4, 2, 1, bias=False),
116             nn.BatchNorm2d(ngf * 4),
117             nn.ReLU(True),
118             # state size. (ngf*4) x 8 x 8
119             nn.ConvTranspose2d(ngf * 4, ngf * 2, 4, 2, 1, bias=False),
120             nn.BatchNorm2d(ngf * 2),
121             nn.ReLU(True),
122             # state size. (ngf*2) x 16 x 16
123             nn.ConvTranspose2d(ngf * 2,      ngf, 4, 2, 1, bias=False),
124             nn.BatchNorm2d(ngf),
125             nn.ReLU(True),
126             # state size. (ngf) x 32 x 32
127             nn.ConvTranspose2d(      ngf,      nc, 4, 2, 1, bias=False),
128             nn.Tanh()
129             # state size. (nc) x 64 x 64
130         )
131
132     def forward(self, input):
133         if isinstance(input.data, torch.cuda.FloatTensor) and self.ngpu > 1:
134             output = nn.parallel.data_parallel(self.main, input, range(self.ngpu))
135         else:
136             output = self.main(input)
137         return output
```



Finally!

```
140 netG = _netG(ngpu)
141 netG.apply(weights_init)
142 if opt.netG != '':
143     netG.load_state_dict(torch.load(opt.netG))
144 print(netG)
145
```



Finally!

```
147 class _netD(nn.Module):
148     def __init__(self, ngpu):
149         super(_netD, self).__init__()
150         self.ngpu = ngpu
151         self.main = nn.Sequential(
152             # input is (nc) x 64 x 64
153             nn.Conv2d(nc, ndf, 4, 2, 1, bias=False),
154             nn.LeakyReLU(0.2, inplace=True),
155             # state size. (ndf) x 32 x 32
156             nn.Conv2d(ndf, ndf * 2, 4, 2, 1, bias=False),
157             nn.BatchNorm2d(ndf * 2),
158             nn.LeakyReLU(0.2, inplace=True),
159             # state size. (ndf*2) x 16 x 16
160             nn.Conv2d(ndf * 2, ndf * 4, 4, 2, 1, bias=False),
161             nn.BatchNorm2d(ndf * 4),
162             nn.LeakyReLU(0.2, inplace=True),
163             # state size. (ndf*4) x 8 x 8
164             nn.Conv2d(ndf * 4, ndf * 8, 4, 2, 1, bias=False),
165             nn.BatchNorm2d(ndf * 8),
166             nn.LeakyReLU(0.2, inplace=True),
167             # state size. (ndf*8) x 4 x 4
168             nn.Conv2d(ndf * 8, 1, 4, 1, 0, bias=False),
169             nn.Sigmoid()
170         )
171
172     def forward(self, input):
173         if isinstance(input.data, torch.cuda.FloatTensor) and self.ngpu > 1:
174             output = nn.parallel.data_parallel(self.main, input, range(self.ngpu))
175         else:
176             output = self.main(input)
177
178         return output.view(-1, 1).squeeze(1)
```



Finally!

```
181 netD = _netD(ngpu)
182 netD.apply(weights_init)
183 if opt.netD != '':
184     netD.load_state_dict(torch.load(opt.netD))
185 print(netD)
186
```



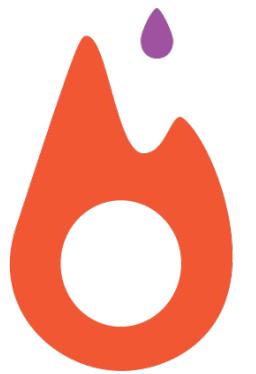
Finally!

```
187     criterion = nn.BCELoss()
```



Finally!

```
189     input = torch.FloatTensor(opt.batchSize, 3, opt.imageSize, opt.imageSize)
190     noise = torch.FloatTensor(opt.batchSize, nz, 1, 1)
191     fixed_noise = torch.FloatTensor(opt.batchSize, nz, 1, 1).normal_(0, 1)
192     label = torch.FloatTensor(opt.batchSize)
193     real_label = 1
194     fake_label = 0
195
196     if opt.cuda:
197         netD.cuda()
198         netG.cuda()
199         criterion.cuda()
200         input, label = input.cuda(), label.cuda()
201         noise, fixed_noise = noise.cuda(), fixed_noise.cuda()
202
203     fixed_noise = Variable(fixed_noise)
```



Finally!

```
205 # setup optimizer  
206 optimizerD = optim.Adam(netD.parameters(), lr=opt.lr, betas=(opt.beta1, 0.999))  
207 optimizerG = optim.Adam(netG.parameters(), lr=opt.lr, betas=(opt.beta1, 0.999))  
208
```



Finally!

```
209 for epoch in range(opt.niter):
210     for i, data in enumerate(dataloader, 0):
211         #####
212         # (1) Update D network: maximize log(D(x)) + log(1 - D(G(z)))
213         #####
214         # train with real
215         netD.zero_grad()
216         real_cpu, _ = data
217         batch_size = real_cpu.size(0)
218         if opt.cuda:
219             real_cpu = real_cpu.cuda()
220             input.resize_as_(real_cpu).copy_(real_cpu)
221             label.resize_(batch_size).fill_(real_label)
222             inputv = Variable(input)
223             labelv = Variable(label)
224
225             output = netD(inputv)
226             errD_real = criterion(output, labelv)
227             errD_real.backward()
228             D_x = output.data.mean()
229
230         # train with fake
231         noise.resize_(batch_size, nz, 1, 1).normal_(0, 1)
232         noisev = Variable(noise)
233         fake = netG(noisev)
234         labelv = Variable(label.fill_(fake_label))
235         output = netD(fake.detach())
236         errD_fake = criterion(output, labelv)
237         errD_fake.backward()
238         D_G_z1 = output.data.mean()
239         errD = errD_real + errD_fake
240         optimizerD.step()
```



Final

```
242 #####  
243 # (2) Update G network: maximize log(D(G(z)))  
244 #####  
245 netG.zero_grad()  
246 labelv = Variable(label.fill_(real_label)) # fake labels are real for generator cost  
247 output = netD(fake)  
248 errG = criterion(output, labelv)  
249 errG.backward()  
250 D_G_z2 = output.data.mean()  
251 optimizerG.step()  
252  
253 print(' [%d/%d] [%d/%d] Loss_D: %.4f Loss_G: %.4f D(x): %.4f D(G(z)): %.4f / %.4f'  
254     % (epoch, opt.niter, i, len(dataloader),  
255         errD.data[0], errG.data[0], D_x, D_G_z1, D_G_z2))  
256 if i % 100 == 0:  
257     vutils.save_image(real_cpu,  
258                     '%s/real_samples.png' % opt.outf,  
259                     normalize=True)  
260     fake = netG(fixed_noise)  
261     vutils.save_image(fake.data,  
262                     '%s/fake_samples_epoch_%03d.png' % (opt.outf, epoch),  
263                     normalize=True)  
264  
265 # do checkpointing  
266 torch.save(netG.state_dict(), '%s/netG_epoch_%d.pth' % (opt.outf, epoch))  
267 torch.save(netD.state_dict(), '%s/netD_epoch_%d.pth' % (opt.outf, epoch))
```



