



A dynamic view of the deep learning world

Soumith Chintala
Facebook AI Research

Overview

- What is Torch?
- The Community
- Today's AI
- What Next?



What is torch ?

- Interactive Scientific computing framework

Strings, numbers, tables - a tiny introduction

```
In [ ]: a = 'hello'  
In [ ]: print(a)  
In [ ]: b = {}  
In [ ]: b[1] = a  
In [ ]: print(b)  
In [ ]: b[2] = 30  
In [ ]: for i=1,#b do -- the # operator is the length operator in Lua  
         print(b[i])  
      end
```



What is torch ?

- Interactive Scientific computing framework

Tensors

```
In [ ]: a = torch.Tensor(5,3) -- construct a 5x3 matrix, uninitialized
```

```
In [ ]: a = torch.rand(5,3)  
print(a)
```

```
In [ ]: b=torch.rand(3,4)
```

```
In [ ]: -- matrix-matrix multiplication: syntax 1  
a*b
```

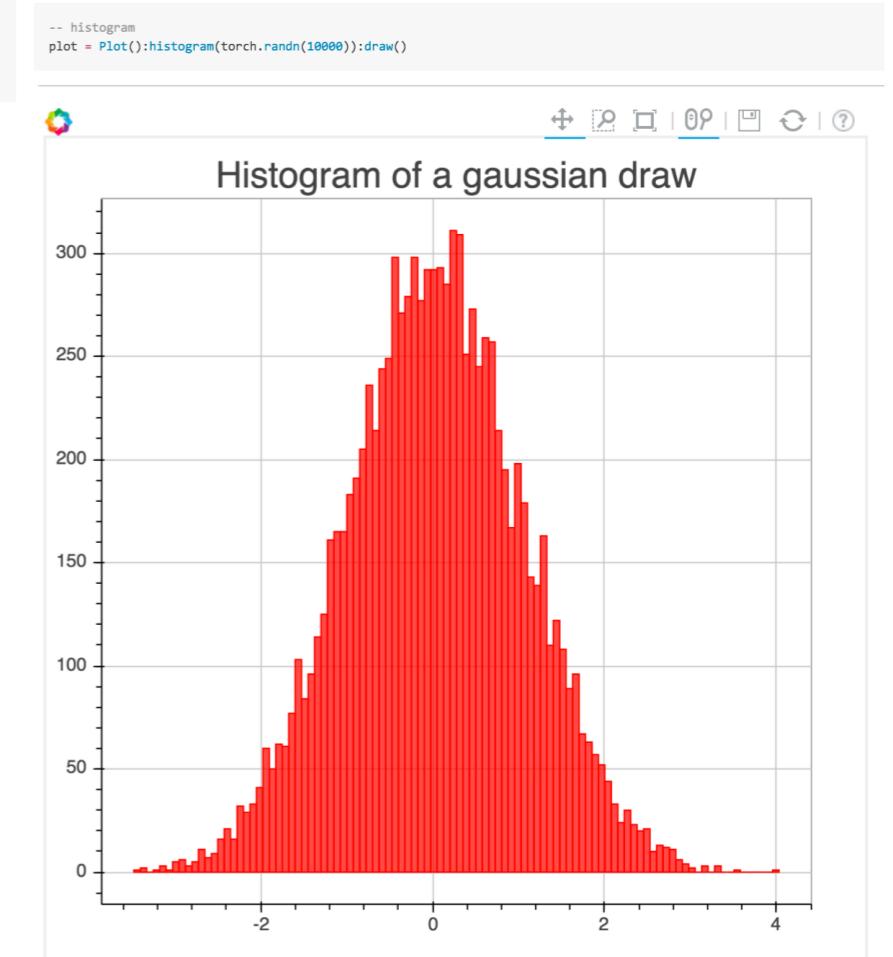
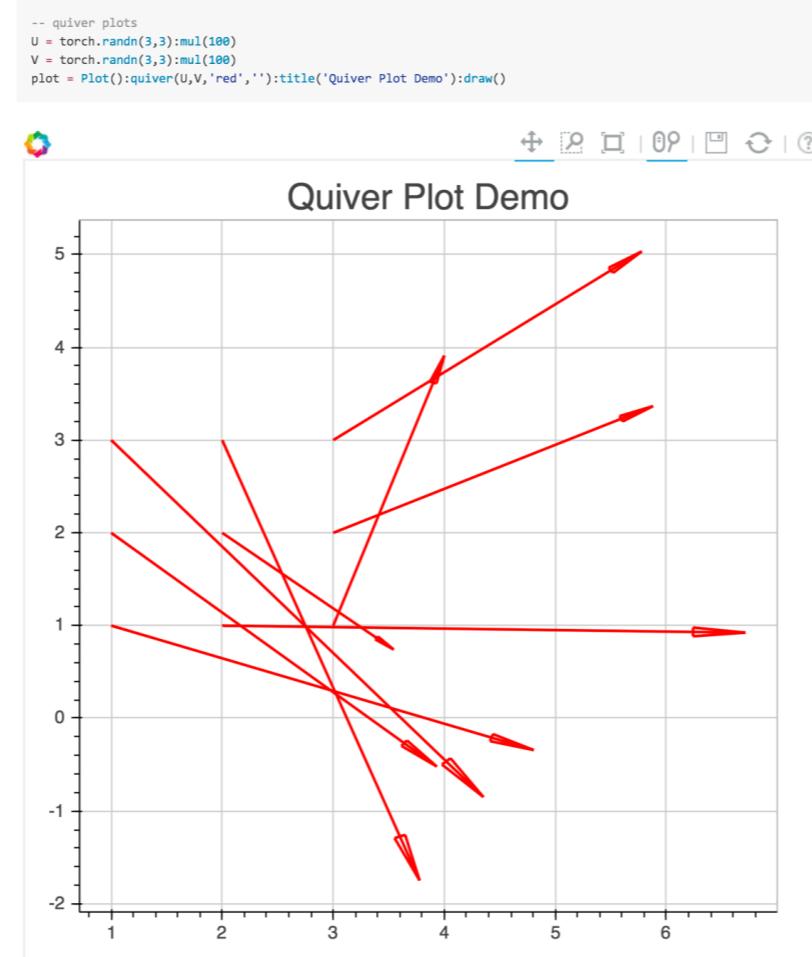
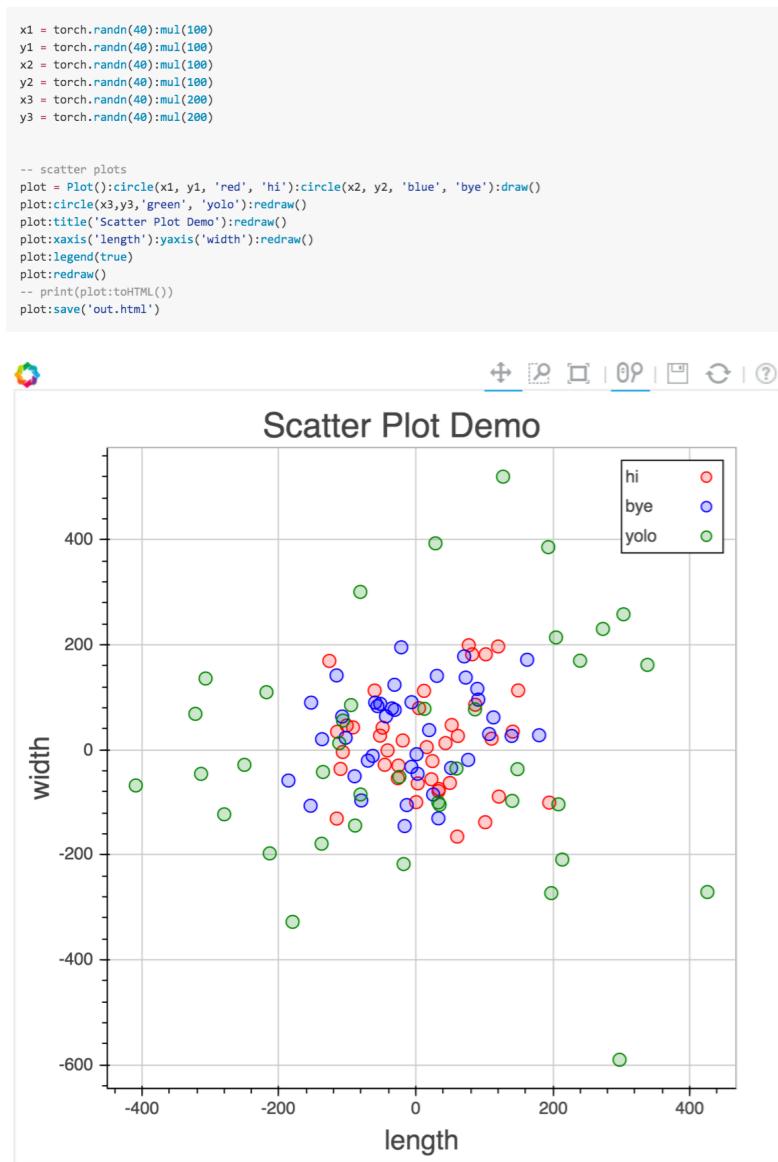
```
In [ ]: -- matrix-matrix multiplication: syntax 2  
torch.mm(a,b)
```

```
In [ ]: -- matrix-matrix multiplication: syntax 3  
c=torch.Tensor(5,4)  
c:mm(a,b) -- store the result of a*b in c
```



What is torch ?

- Similar to Matlab / Python+Numpy



What is torch ?

- Easy integration into and from C
- Example: using CuDNN functions

```
for g = 0, self.groups - 1 do
    errcheck('cudnnConvolutionForward', cudnn.getHandle(),
              one:data(),
              self.iDesc[0], input:data() + g*self.input_offset,
              self.weightDesc[0], self.weight:data() + g*self.weight_offset,
              self.convDesc[0], self.fwdAlgType[0],
              self.extraBuffer:data(), self.extraBufferSizeInBytes,
              zero:data(),
              self.oDesc[0], self.output:data() + g*self.output_offset);
end
```



What is torch ?

- Strong GPU support

CUDA Tensors

Tensors can be moved onto GPU using the :cuda function

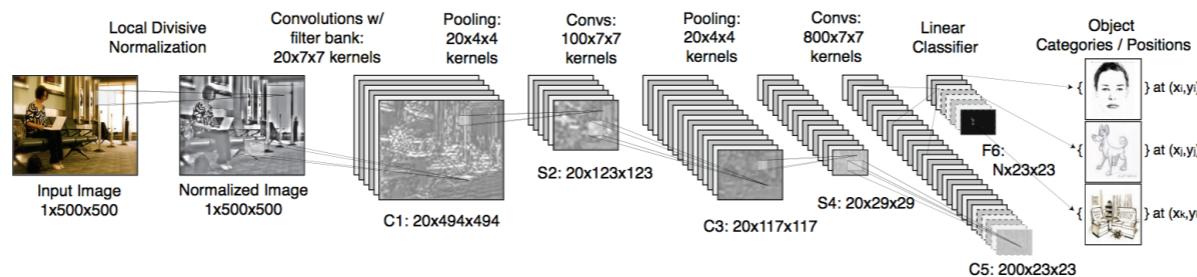
```
In [ ]: require 'cutorch';
          a = a:cuda()
          b = b:cuda()
          c = c:cuda()
          c:mm(a,b) -- done on GPU
```



Neural Networks

- nn: neural networks made easy
- building blocks of differentiable modules
 - **define a model with pre-normalization, to work on raw RGB images:**

```
01 model = nn.Sequential()  
02  
03 model.add(nn.SpatialConvolution(3,16,5,5))  
04 model.add(nn.Tanh())  
05 model.add(nn.SpatialMaxPooling(2,2,2,2))  
06 model.add(nn.SpatialContrastiveNormalization(16, image.gaussian(3)))  
07  
08 model.add(nn.SpatialConvolution(16,64,5,5))  
09 model.add(nn.Tanh())  
10 model.add(nn.SpatialMaxPooling(2,2,2,2))  
11 model.add(nn.SpatialContrastiveNormalization(64, image.gaussian(3)))  
12  
13 model.add(nn.SpatialConvolution(64,256,5,5))  
14 model.add(nn.Tanh())  
15 model.add(nn.Reshape(256))  
16 model.add(nn.Linear(256,10))  
17 model.add(nn.LogSoftMax())
```



autograd by

- Write imperative programs
- Backprop defined for every operation in the language

```
neuralNet = function(params, x, y)
    local h1 = t.tanh(x * params.W[1] + params.b[1])
    local h2 = t.tanh(h1 * params.W[2] + params.b[2])
    local yHat = h2 - t.log(t.sum(t.exp(h2)))
    local loss = - t.sum(t.cmul(yHat, y))
    return loss
end

-- gradients:
dneuralNet = grad(neuralNet)

-- some data:
x = t.randn(1,100)
y = t.Tensor(1,10):zero() y[1][3] = 1

-- compute loss and gradients wrt all parameters in params:
dparams, loss = dneuralNet(params, x, y)
```



Distributed Learning

- Multi-GPU and Multi-Node
- distlearn by 
- MPI-like model
- scales to a large amount of nodes

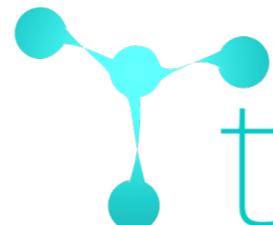
```
-- Use a tau of 10 and an alpha of 0.2
local allReduceEA = require 'distlearn.AllReduceEA'(tree, 10, 0.2)
-- Make sure all the nodes start with the same parameter values
allReduceEA.synchronizeParameters(params)
for _ = 1,epochs do
    for _ = 1,steps
        -- Compute your gradients as normal
        local grads = computeYourGrads(...)
        -- Do your SGD as normal
        SGD(params, grads)
        -- Average the params
        allReduceEA.averageParameters(params)
    end
    -- Make sure the center's haven't drifted too far due to
    -- floating point precision error build up
    allReduceEA.synchronizeCenter(params)
    -- Validate...
end
```



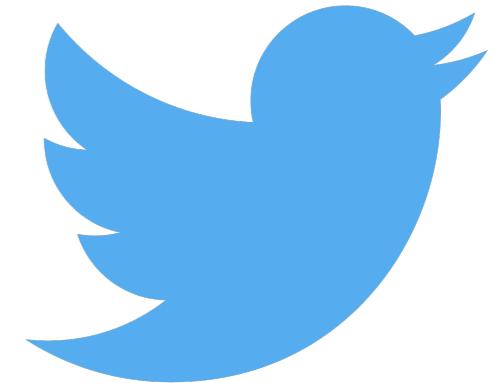
Core Philosophy

- Interactive computing
 - No compilation time
- Imperative programming
 - Easy to understand, think and debug
- Minimal abstraction
 - Thinking linearly
- Maximal Flexibility
 - No constraints on interfaces or classes





torch Community



canary

Inria



PURDUE
UNIVERSITY

Yandex

element™



AMD

IBM®

MULTICORE
WARE





[szagoruyko / loadcaffe](#)

Watch ▾ 19 Unstar 203 Fork 69

Code

Issues 10

Pull requests 1

Wiki

Pulse

Graphs

Load Caffe networks in Torch7

[facebook / fb.resnet.torch](#)

Unwatch ▾ 62 Unstar 421 Fork 85

Code

Issues 4

Pull requests 0

Wiki

Pulse

Graphs

Torch implementation of ResNet from <http://arxiv.org/abs/1512.03385> and training scripts

[Moodstocks / inception-v3.torch](#)

Watch ▾ 11 Unstar 48 Fork 8

Code

Issues 1

Pull requests 0

Wiki

Pulse

Graphs

Rethinking the Inception Architecture for Computer Vision <http://arxiv.org/abs/1512.00567>





karpathy / [neuraltalk2](#)

Watch 171 Unstar 2,153 Fork 320

Code Issues 45 Pull requests 4 Wiki Pulse Graphs

Efficient Image Captioning code in Torch, runs on GPU

NeuralTalk2

Recurrent Neural Network captions your images. Now much faster and better than the original [NeuralTalk](#). Compared to the original NeuralTalk this implementation is **batched, uses Torch, runs on a GPU, and supports CNN finetuning**. All of these together result in quite a large increase in training speed for the Language Model (~100x), but overall not as much because we also have to forward a VGGNet. However, overall very good models can be trained in 2-3 days, and they show a much better performance.

This is an early code release that works great but is slightly hastily released and probably requires some code reading of inline comments (which I tried to be quite good with in general). I will be improving it over time but wanted to push the code out there because I promised it to too many people.

This current code (and the pretrained model) gets ~0.9 CIDEr, which would place it around spot #8 on the [codalab leaderboard](#). I will submit the actual result soon.



a man is playing tennis on a tennis court



a train is traveling down the tracks at a train station



a cake with a slice cut out of it



a bench sitting on a patch of grass next to a sidewalk

You can find a few more example results on the [demo page](#). These results will improve a bit more once the last few bells and whistles are in place (e.g. beam search, ensembling, reranking).

There's also a fun video by [@kcimc](#), where he runs a neuraltalk2 pretrained model in real time on his laptop during a walk in Amsterdam.





jcjohnson / **neural-style** Watch ▾ 389 Unstar 7,152 Fork 901

Code Issues 98 Pull requests 12 Wiki Pulse Graphs

Torch implementation of neural style algorithm

neural-style

This is a torch implementation of the paper [A Neural Algorithm of Artistic Style](#) by Leon A. Gatys, Alexander S. Ecker, and Matthias Bethge.

The paper presents an algorithm for combining the content of one image with the style of another image using convolutional neural networks. Here's an example that maps the artistic style of [The Starry Night](#) onto a night-time photograph of the Stanford campus:





OpenNMT / OpenNMT

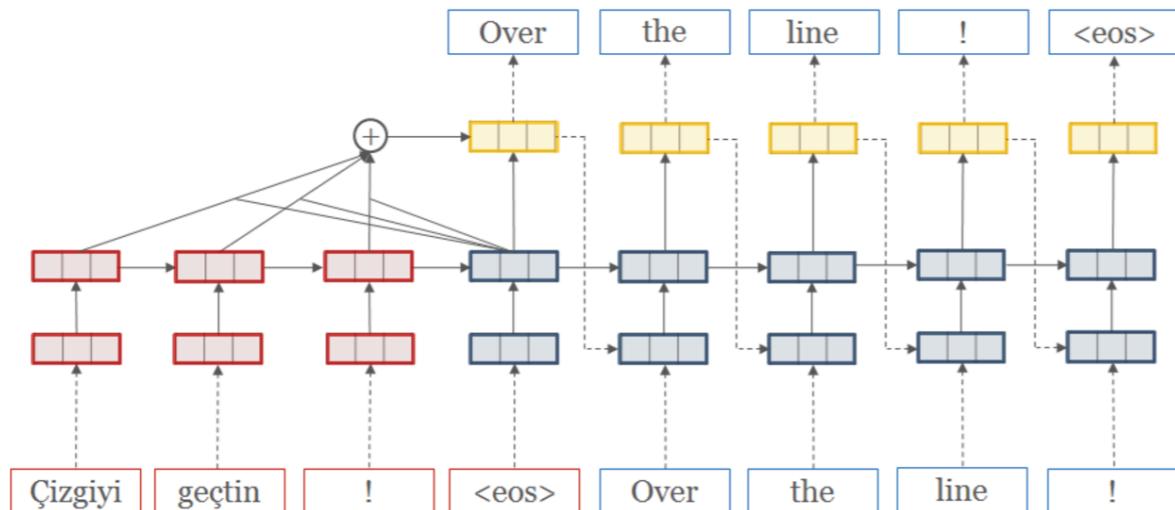
Watch ▾ 64 Unstar 601 Fork 121

Code Issues 8 Pull requests 4 Projects 2 Wiki Pulse Graphs

Open-Source Neural Machine Translation in Torch <http://opennmt.net/>

OpenNMT: Open-Source Neural Machine Translation

OpenNMT is a full-featured, open-source (MIT) neural machine translation system utilizing the Torch mathematical toolkit.



The system is designed to be simple to use and easy to extend , while maintaining efficiency and state-of-the-art translation accuracy. Features include:

- Speed and memory optimizations for high-performance GPU training.
- Simple general-purpose interface, only requires source/target data files.
- [C++ implementation of the translator](#) for easy deployment.
- Extensions to allow other sequence generation tasks such as summarization and image captioning.



torch Community

DmitryUlyanov / **fast-neural-doodle**

Code Issues Pull requests Wiki Pulse Graphs

Watch 3 Unstar 111 Fork 13

Faster neural doodle



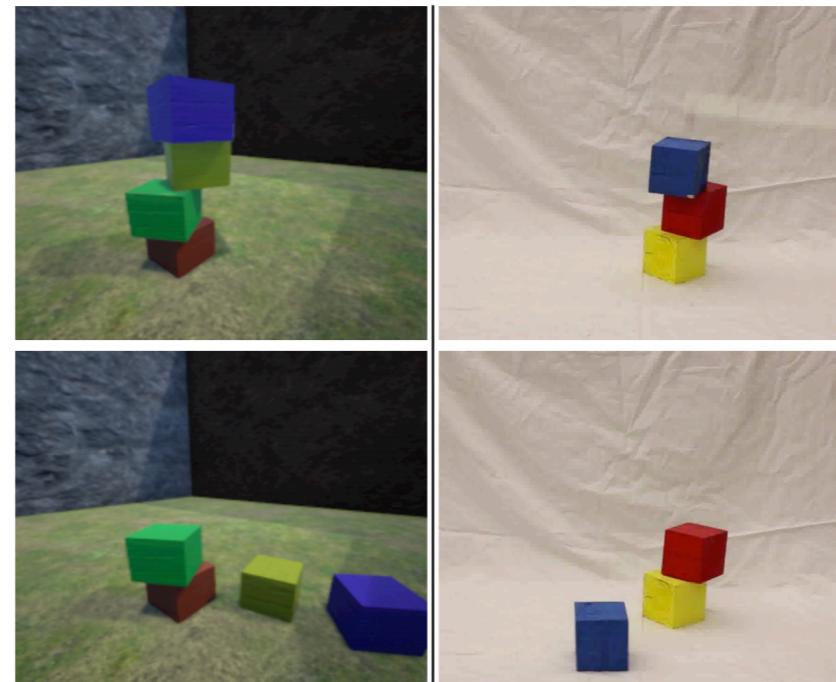
torch Community

facebook / UETorch

Code Issues Pull requests Wiki Pulse Graphs

Unwatch 37 Unstar 135 Fork 20

A Torch plugin for Unreal Engine 4.



TorchCraft / TorchCraft

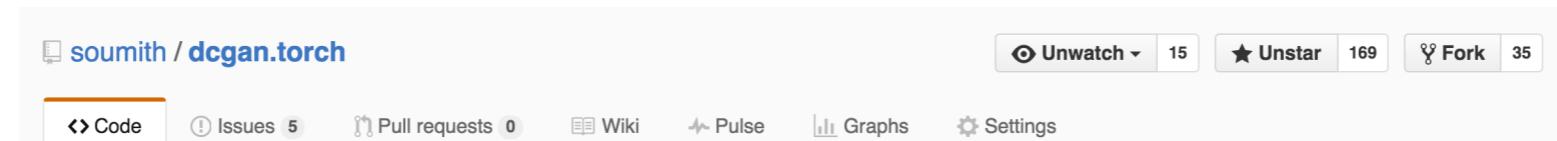
Code Issues 4 Pull requests 0 Projects 0 Wiki Pulse Graphs Settings

Unwatch 43 Unstar 608 Fork 77

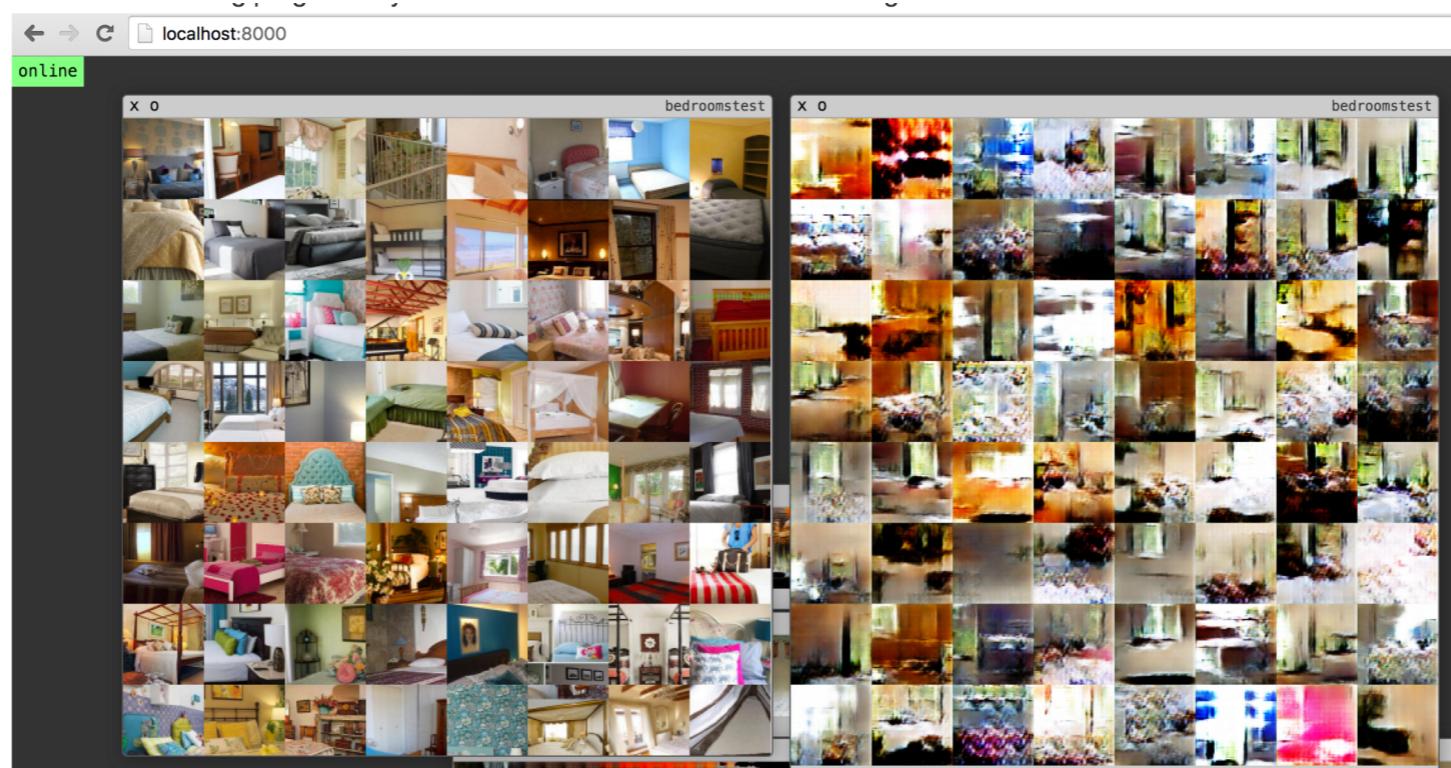
Connecting Torch to StarCraft

Edit lurch

torch Community



A torch implementation of <http://arxiv.org/abs/1511.06434> — Edit



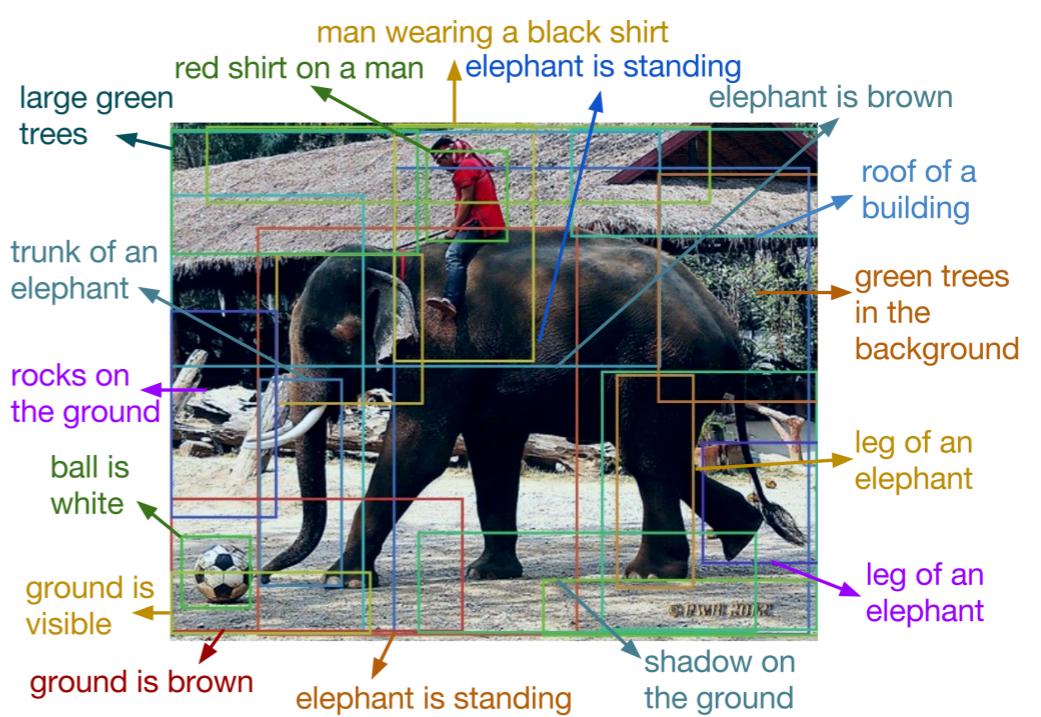
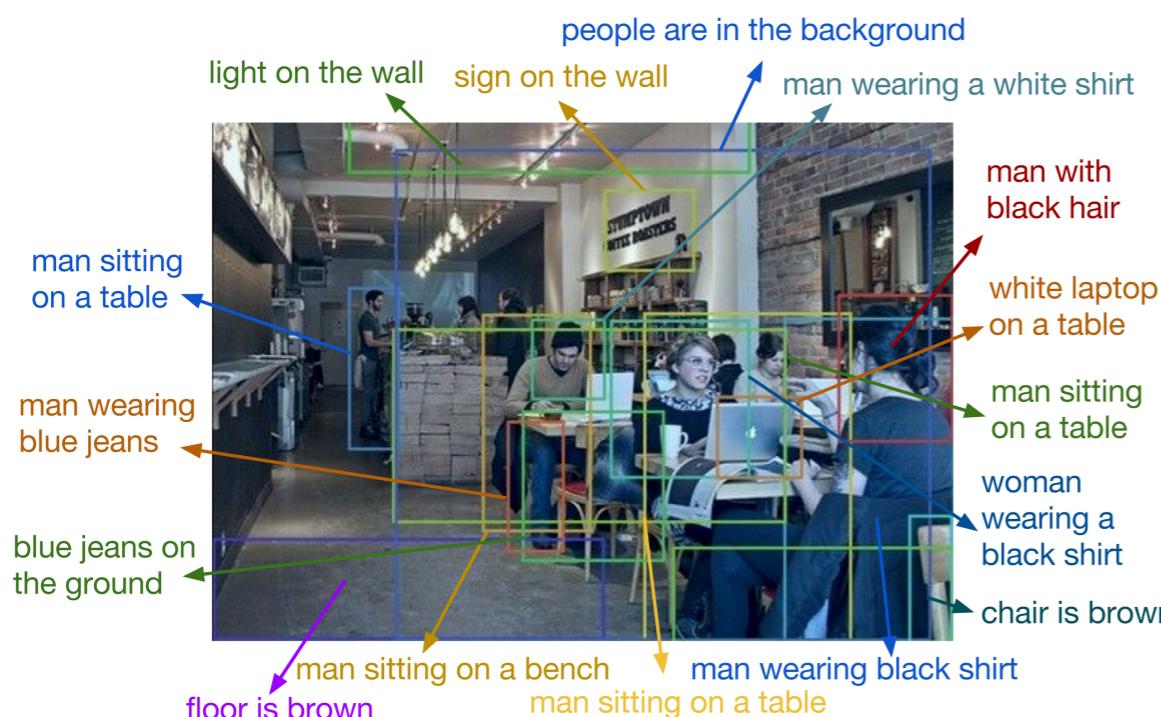
Today's AI

- Reasonably solved to be used in production



Today's AI

- Classification and Detection
- in images, videos, volumetric data, sparse datasets



DenseCap: Johnson et. al. <https://github.com/jcjohnson/densecap>



Today's AI

- Segmentation



DeepMask: Pinhero et. al.

Today's AI

- Text Classification (sentiment analysis etc.)
- Text Embeddings
- Graph embeddings
- Machine Translation
- Ads ranking



Today's AI

static datasets + static model structure



Today's AI

static datasets + static model structure

model does not change over training time



Today's AI

static datasets + static model structure

model does not change over training time

offline learning



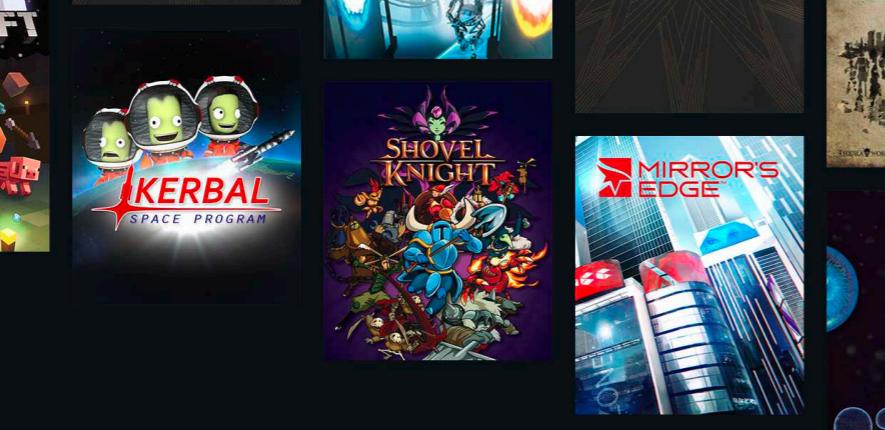
Active and Future AI Research

- Agents training in different and new environments



Active and Future AI Research

- RL Agents in different and new environments
 - Example: OpenAI Universe



1000+ environments

Universe contains an ever-expanding catalog of environments. Environments will soon include programs from EA, Microsoft Studios, Valve, Zachtronics, and others.

[View on GitHub >](#)

Human-like interface

Agents use the same senses and controls as humans: seeing pixels and using a keyboard and mouse. Universe makes it possible to train a single agent on any task a human can complete with a computer.

eye + ⌂ + keyboard

Active and Future AI Research

- RL Agents in different and new environments
 - Example: OpenAI Universe
- Online Learning
 - Example: self-driving cars



Active and Future AI Research

- RL Agents in different and new environments
 - Example: OpenAI Universe
- Online Learning
 - Example: self-driving cars
- Dynamic Neural Networks
 - self-adding new memory or layers
 - changing evaluation path based on inputs



Active and Future AI Research

- RL Agents in different and new environments
 - Example: OpenAI Universe
- Online Learning
 - Example: self-driving cars
- Dynamic Neural Networks
 - self-adding new memory or layers
 - changing evaluation path based on inputs
- Structured Prediction
 - Viterbi style decoders



A Next-Gen Framework

- Interop with environments
 - Like Universe, VizDoom, etc.
 - Easy data loaders
- Rich scientific ecosystem (such as SciPy)
 - Neural Networks are not the only methods to be used
- Dynamic Neural Networks
 - with no specific or static structure
- Minimal Abstractions
 - better debugging and low-level programming in the hands of researcher
- Graph Compilers
 - To speed up structured prediction and fuse operations
 - Numba, XLA, Theano



Let's share!

